

Apache Tika Documentation

Table of Contents

- [Where to start](#)

This reference guide was generated with the assistance of AI and requires human review before it can be fully trusted. This documentation serves as an example and a starting point, but more work remains. Contributions and corrections are welcome.

The Apache Tika toolkit detects and extracts metadata and text from over a thousand different file types (such as PPT, XLS, and PDF). All of these file types can be parsed through a single interface, making Tika useful for search engine indexing, content analysis, translation, and much more. Tika identifies a file's type from its bytes, so a missing or misleading extension does not change the result.

Where to start

- [Supported Formats](#) — what Tika parses, family by family, plus the generated parser-to-media-type list.
- [Using Tika](#) — choose an integration method: Java API, command line, REST server, or gRPC.
- [Pipes](#) — process-isolated, fault-tolerant extraction at scale.
- [Configuration](#) — the JSON config file: parsers, detectors, metadata filters.
- [Setting Limits](#) and [Robustness](#) — bounding time, memory and output when parsing content you do not control.
- [Security](#) — the trust model, and what to lock down before exposing tika-server or tika-grpc.
- [Migrating to 4.x](#) — what changed since 3.x and what you have to update.
- [FAQ](#) — the questions upgraders ask most.

Apache Tika is an Apache Software Foundation project, formerly a subproject of Apache Lucene.

Built from commit: 47c1a0335 (2026-09-14)